

## **Predictive Analytics Case Study**

Surenther Selvaraj

Bellevue University

DSC630: Predictive Analytics

Professor: Frank Neugebauer

Date: Jan 23, 2026

# 1. Introduction

## The Problem Being Solved

At its core, Netflix faced a "search and discovery" problem. In the early 2000s, as its library transitioned from physical DVDs to streaming, the sheer volume of content became overwhelming for the average user. The specific problem was customer churn caused by "choice paralysis." If a user spent 60 to 90 seconds or more looking for something to watch without success, the likelihood of them closing the application and canceling their subscription increased significantly. Netflix needed a way to predict, with high accuracy, what specific title a user would enjoy watching right now.

## Importance of the Problem

For a subscription-based model, retention is the primary driver of revenue. Unlike traditional TV, where viewers are "pushed" content based on a schedule, Netflix relies on a "pull" model. If the recommendation engine fails, the service's value proposition diminishes. By 2024 standards, it is estimated that Netflix's recommendation engine is responsible for over 80% of the content watched on the platform, saving the company over \$1 billion annually in customer retention costs.

## Data Acquisition

Netflix acquires data through a multi-channel pipeline that captures both **explicit** and **implicit** signals:

- **Explicit Data:** Ratings (the old 1–5-star system, Thumbs Up/Down), user reviews, and "My List" additions.
- **Implicit Data:** This is the most critical volume of data. It includes play duration, time of day watched, device type, browsing behavior (what they hovered over but didn't click), and even "rewatch" patterns.
- **Metadata:** Netflix employs thousands of people to manually tag content with "micro-genres," moods, and character archetypes, creating a rich dataset of content features to pair with user behavior.

## 2. Methods and Results

### Data Preparation Steps

Before modeling could begin, the raw logs of billions of stream events had to be cleaned and structured. Key steps included:

1. **Normalization:** Adjusting for "rating bias." Some users are "easy graders" (averaging 4 stars), while others are "harsh" (averaging 2 stars). The data was centered to reflect relative preferences.
2. **Handling Sparsity:** Most users have only watched a tiny fraction of the library. The Netflix team used matrix completion techniques to fill in the "holes" in the user-item matrix.
3. **Time-Decay Filtering:** Older data was weighted less heavily than recent data, as user tastes evolve over time

### How the Problem Was Solved: Modeling Techniques

Netflix moved away from a single algorithm to an Ensemble Method, which combines multiple models to produce a single, more accurate prediction.

The primary techniques used include:

- **Collaborative Filtering (Matrix Factorization):** This looks at patterns of users who are similar to you. If User A and User B both liked *Stranger Things*, and User B liked *Dark*, the model predicts User A will also like *Dark*.
- **Content-Based Filtering:** This focuses on the properties of the items. If you watch many Sci-Fi movies directed by Christopher Nolan, the system recommends other Sci-Fi movies or Nolan films.
- **Deep Learning (Neural Networks):** Specifically, Restricted Boltzmann Machines (RBMs) were used to capture non-linear relationships in the data that simple linear models missed.

- **Multi-Armed Bandits:** used for "Explore vs. Exploit." The system occasionally shows you something completely outside your comfort zone to see if your tastes have changed, preventing the "filter bubble" effect.

## Why These Methods?

Netflix chose ensemble methods because "no single model wins." While Matrix Factorization is excellent at finding broad similarities, it fails with "Cold Starts" (new movies with no history). Content-based filtering solves the Cold Start problem. By layering these models, Netflix ensures that every row on your homepage (e.g., "Trending Now," "Because You Watched...") is powered by a different specialized sub-model.

## Evaluation Metrics

The primary metric during the initial development was Root Mean Square Error (RMSE).

- **Why RMSE?** It penalizes large errors more heavily than small ones. In a recommendation context, predicting a 5-star movie as a 1-star (a "false negative") is a massive failure that ruins user trust, so RMSE was used to minimize those "catastrophic" mis-predictions.
- **Modern Metrics:** Today, Netflix uses Precision at K (did the user click on one of the top 10 items shown?) and Take-Rate, which measures the actual consumption of recommended titles.

## 3. Conclusion

### Implementation of the Model

The model is implemented via a distributed cloud architecture (AWS). It is not a static batch process; it is a Real-Time Streaming system. As you finish an episode, the "Next Episode" or "Recommended for You" list updates instantly based on your most recent interaction. Furthermore, Netflix implements Artwork Personalization: the movie thumbnails you see are tailored to you. If you like romances, you might see the lead couple; if you like action, you might see a car chase from the same movie.

### Actionable Consequences

The most significant consequence was the shift from Licensed Content to Original Content. By analyzing their predictive data, Netflix realized there was a massive,

underserved demand for "Political Dramas" and "David Fincher fans." This led to the \$100 million investment in House of Cards, their first major original hit. They didn't "guess" it would be a hit; their predictive models told them the audience already existed.

## Lessons Learned

Netflix learned that **context is king**. A user's preference on a Friday night (looking for a blockbuster) is different from their preference on a Monday morning (looking for a 20-minute sitcom). This led to the development of "contextual bandits" that factor in the day of the week and the device being used.

## Future Improvements and Different Approaches

If Netflix were to start over today, they would place less emphasis on explicit ratings (stars) and more on Reinforcement Learning (RL).

- **Why RL?** Traditional models are "greedy"—they try to get you to click now. Reinforcement Learning looks at the "Long-Term Value" of a user session. Instead of recommending a "clickbait" movie that you turn off after 5 minutes, an RL model might recommend a slower-paced series that keeps you subscribed for the next three months.
- **Causal Inference:** Netflix is also moving toward understanding why a user chose a show, not just that they chose it, to better distinguish between a user's genuine interest and the "position bias" of where the show was placed on the screen.

## References

**Amatriain, X., & Basilico, J. (2012).** Netflix recommendations: Beyond the 5 stars (Part 1). *The Netflix Tech Blog*.

<https://medium.com/netflix-techblog/netflix-recommendations-beyond-the-5-stars-part-1-55838468f429>

**Bennett, J., & Lanning, S. (2007).** The Netflix Prize. *Proceedings of KDD Cup and Workshop*.

<https://www.cs.uic.edu/~liub/KDD-cup-2007/proceedings/The-Netflix-Prize-Bennett.pdf>

**Chandrashekar, A., Amat, F., Basilico, J., & Jebara, T. (2017).** Artwork personalization at Netflix. *The Netflix Tech Blog*.

<https://netflixtechblog.com/artwork-personalization-at-netflix-ed51144432f7>

**Gomez-Uribe, C. A., & Hunt, N. (2015).** The Netflix recommender system: Algorithms, business value, and innovation. *ACM Transactions on Management Information Systems (TMIS)*, 6(4), 1–19.

<https://dl.acm.org/doi/10.1145/2843948>

**Hallinan, B., & Striphas, T. (2016).** Recommended for you: The Netflix Prize and the production of algorithmic culture. *New Media & Society*, 18(1), 117–137.

<https://journals.sagepub.com/doi/abs/10.1177/1461444814538646>.